PM – Programming Parallel Models

Version 0.1

Language Reference (incomplete)

© Tim Bellerby, University of Hull, UK

Contents

Background

PM (Parallel Models) is a new programming language designed for implementing environmental models, particularly in a research context where ease of coding and performance of the implementation are both essential but frequently conflicting requirements. PM is designed to allow a natural coding style, including familiar forms such as loops over array elements. Language elements have been included (or more importantly excluded) from the language in a careful combination that enables the language implementation to both vectorise and parallelise the code. The following goals guided the development of the PM language:

- Language structures facilitate the generation of vectorised and parallelised code
- Vectorisation and parallelisation should be explicit and configurable
- Programmers should be able to concentrate on coding the algorithm – not on how to get it to run efficiently using whole-array operations and/or in parallel
- Support for data structures required in modelling: matrices, vectors, grids (including grids of matrices, vectors and sub-grids) and meshes with generalised connectivity.
- PM code should be as readable as possible by those not familiar with the language.
- PM should be formally specified – not defined by an implementation.

PM draws inspiration from many sources. Notable influences are: Parasail (parasail-lang.org), Go (golang.org) , NESL (http://www.cs.cmu.edu/~scandal/nesl.html) and llc (Reyes et al.,2009, 16[th] European PVM/MPI Users Group Meeting). Note should also be made of two recent languages tackling similar problems: Chapel (chapel.cray.com) and Julia (julialang.org).

Conventions used to define language syntax

PM syntax will be described using the following extended BNF notation:

name ::= *list*	Define a non-terminal element in terms of other elements
name	Non terminal element
dim	Keyword
'>='	Character combination
[*elements*]	Elements are optional – may appear zero or one times
{ *elements* }	Elements may appear zero, one or more times
el1 | *el2*	Exactly one of the listed element sequences must be present
(el1 | el2)	Brackets may be used to group selections that are not enclosed by {} or []

Lexical Elements

Comments start with a '!' and continue to the end of the line

```
! Comments are ignored
```

PM uses upper and lower case letters, digits and a number of special symbols:

letter::= 'a' \|'b' \|'c' \|'d' \|'e' \|'f' \|'g' \|'h' \|'i' \|'j' \|'k' \|'l' \|'m' \|'n' \|'o' \|'p' \|'q' \|'r' \|'s' \|'t' \|'u' \|'v' \|'w' \|'x' \|'y' \| 'z'
  'A' \|'B' \|'C' \|'D' \|'E' \|'F' \|'G' \|'H' \|'I' \|'J' \|'K' \|'L' \|'M' \|'N' \|'O' \|'P' \|'Q' \|'R' \|'S' \|'T' \|'U' \|'V' \|'W' \|'X' \|'Y' \| 'Z'
digit::= '0' \| '1' \| '2' \|'3' \| '4' \| '5' \| '6' \| '7' \| '8' \| '9'

The following single characters and character combinations are used as ***delimiters***:

'$'	'&'	'('	')'	'**'	'*'	','	'//'	'/'	'=>'	':'	'||'	';'	'@'	'['
']'	'_'	'{'	'|'	'}'	'-'	'""'	'+'	'<'	'<='	'='	'=='	'>'	'>='	'%'
'..'	'.'	'->'	'::'											

To cater for environments with restricted character sets, the following substitutions are always permitted:

'(/' for '['	'/)' for ']'	'(:' for '{'	':)' for '}'	'%%' for '@'	':/' for '|'
'://' for '||'					

White space characters (space, newline, form feed and horizontal tab) may appear between lexical elements (delimiters, names, keywords, numeric constants) but not within them. Names, keywords and numbers must be separated from each other by white space. Otherwise white space is optional.

New lines are not generally significant, but in any context where a semicolon '**;**' is expected/allowed by the syntax, the semicolon may be omitted if the following lexical (non-white-space) symbol starts on a new line.

Names have a maximum length of 100 characters. They are comprised of upper and lower case letters, decimal digits, and the underscore character '**_**'. They may not start with a digit. Letter case is significant.

name   ::=   ['**_**']*letter*{'**_**'|*letter*|*digit*}

Names that start with '**_**' are unique to the module within which they are defined. A module entity defined using such a name may not be referred to from another module.

The following are **keywords** and may not be used as names:

also	and	arg	argc	by	case
check	const	debug	default	dim	div
do	else	elseif	enddebug	enddo	endfind
endfor	endif	endlet	endproc	endselect	endtype
endwhile	false	find	follow	for	found
from	high	if	in	include	includes
is	key	let	low	mod	not
null	of	opt	otherwise	over	param
proc	prod	rec	or	reduce	render
repeat	result	select	seq	step	struct
then	true	type	until	use	using
when	with	where	while		

Numeric constants take the following forms*:*

number	::=	[*integer_constant* | *real_constant* | *imaginary_constant*] [*bits*]
integer_constant	::=	[{ *digit* } '**r**'] { [*digit* | *letter*] } ['**l**']
real_constant	::=	{ *digit* } ['**.**'*digit* { *digit* }] ['**e**'['**+**' | '**-**'] *digit* { *digit* }]['**d**']
imaginary_constant	::=	[*integer_constant* | *real_constant*] '**i**'
bits	::=	'**_**' { *digit* }

String constants are enclosed by double quotes'**"**'and may include any character other than '**"**'. They may not span more than one line.

```
"Hello World"
```

Modular Structure

A PM program consists of one or more source files, each corresponding to a *module*. A module defines a collection of *procedures*, *types* and *parameters* (module level constants). A *program module* also contains executable statements. Module names are constructed from a sequence of standard PM names optionally joined by '**.**' delimiters. They are linked to source file names in an implementation-dependant manner. Module names starting with '**lib.**' refer to standard libraries.

```
module::=program_module | library_module
program_module::=
        [ decls ] statements
library_module::=
        decls [ debug statements enddebug ]
modname::=
        name { '.' name }
```

Names may optionally start with an underscore character. Such names are local to the module in which they are used and do not match names in any other module into which they are imported.

```
decls::=
        { import ';' }( import | decl ) { ';'decl }
import::=
        use modname [ modifiers ] { ',' modname [ modifiers ] }
```

Modules may include other modules using a **use** statement. In its simplest form, **use** *module_name* the statement imports all definitions in the named module into the current module. Type and parameter declarations should not conflict with other definitions in the module containing the import statement – whether directly defined or imported from other modules. Imported procedure definitions merge with definitions imported to or declared in the module containing the include statement. This merging joins definitions across a PM program – if two different modules define a procedure with the same name, the merging of their definitions affects both the module imported to and both imported modules.

```
modifiers::=
        '{' modifier { ',' modifier } '}'
modifier::=
        [ type | param | proc ] name '=>' name
```

Name clashes may be avoided by using a qualifying clause in the **use** statement. This lists specific elements to import and optionally renames them:

```
        use model4 {
            type model_params
            param theta =>model_theta
            test =>test_model
        }
```

The optional => operator in this context renames the given element. If an element in the list is not qualified as **type**, **param**, or **proc** then it is assumed to be a procedure (**proc**).

Types, procedures and parameters occupy separate name spaces – they may have the same name as each other.

Type, procedure and parameter declarations follow any **include** statements in the module. The simplest of these is the **param** declaration which simply defines a named constant:

---

paramdec::=
      param *name* '=' *xexpr*

For example:

```
param pi = 3.14159265
```

param statements may refer to each other, irrespective of the order in which they are defined. However, such references must not be recursive or mutually recursive. There is no explicit restriction on the procedures employed within a **param** declaration expression – they may be interpreted as zero-parameter procedures with their own separate name space.

Type and procedure declarations are described later in this document.

In a program module, the optional declarations are followed by executable statements. Such a module may be run as a program. It may not, however, be used by other modules. Library modules may be used by other modules. They may not contain executable statements after the declarations except for an optional single debug statement that should contain module testing code.

Types, values and objects

PM programs operate on *values* stored in *objects*. A *type* consists of a (possibly infinite) set of values and/or other types and is defined using a *type constraint expression*. A type may consist of a set containing only a single value. However, it remains distinct from that value. A *conformance* relationship is defined between types and between values and types:

- A value V conforms to an type T if V$\in$T or V$\in$U and U conforms to T
- Type T conforms to type U if T$\in$U or T$\in$V and V conforms to U

While the PM type system provides flexibility similar to some dynamically typed languages, it is designed for static analysis. Run time type inference and procedure selection should only be necessary in situations where specifically polymorphic values are generated (using **polymorphic types** or **reference objects**). This may require type analysis between and across modules.

The universal type, which includes all other types, is denoted using the keyword **any**. Explicit use of the universal type is usually not required; the absence of a type constraint serves a similar purpose in most cases where a type expression is expected.

User defined types

User defined types associate a name with a given set of type constraint expressions, creating a new type that is the union of the listed types. A type is defined in a type definition:

---

typedecl::=
      type *name* ['**{**' *typeparams* '**}**'] [**in** *namelist*] (**is** *typelist* | **includes** *typelist* | **also includes** *typelist*)
typeparams::=
      name['**:**' *type*] { '**,**' *name*['**:**' *type*] }

The simplest type definition simply associates a name with a set of types:

```
type x is int, struct{x:int,y:int}
```

A *parameterised type* declaration defines a template that may be parameterised by one of more types in order to generate an *actual type* that may be used in constraint expression. An actual type is derived from a type template by providing the correct number of type constraint expressions as type arguments.

```
type point{t:num} is struct{ x:t, y:t}
type integer_point is point{int}
type num_point is point{}
```

If *parameter type constraints* are present, then the type arguments used to create an actual type must conform, parameter by parameter, to the corresponding type parameter type constraints, or be missing. Missing type arguments are assumed to be equal to the corresponding type parameter constraint or to **any** if no such constraint is present.

Type templates are characterised by their name and number of parameters, but not their respective type parameter constraints. Thus **R{t}**, **T{t}** and **T{t,u}** each refer to different type templates while **T{t,u}**, **T{r,s}** and **T{,}** refer to the same type template . Two type template definitions may not use the same template name if they refer to the same number of parameters.

An *open type declaration* using the **includes** syntax provides an incomplete definition for a type which may be added to by other definitions. The type definition may be augmented using an **also includes** type definition or by using by declaring another expression to be **in** that type:

```
type a includes int, real
type a also includes string
type b in a is cpx     ! Type a is int, real, string, cpx
```

Augmenting type declarations do not have to be in the same module as the original **includes** declaration and do not have to precede the **includes** declaration in the module or program text.

Open type declarations may also be parameterised:

```
type point{T:num} includes rec point{ x:T,y:T }
type point{T} also includes rec point{ x:T,y:T,z:T}
```

If an **also includes** type definition is parameterised, then the type constraints associated with its parameters must either be absent or must conform, parameter by parameter, with the type constraints in the **includes** definition to which it refers. If type constraints are present, the **also includes** definition only augments the type template when the type arguments in the actual type expression conform to the types parameters in the **also includes** definition. For example:

```
 type point{T:int16} also includes rec point{ combined_index: int32 }

 proc f (  t : point{int},    ! Matches rec point{x:int,y:int}
           u : point{int16)   ! Matches both rec point{x:int,y:int} and
                              !  rec point{ combined_index:int32 }
       )  ...
```

A type definition may not refer to itself recursively unless that recursive reference occurs within the definition of an *optional type*.

The PM system defines a number of **intrinsic types**. These are in effect user defined types, declared in a system module implicitly imported into every other module.

<u>Numeric types</u>

PM supports a range of integer types, defined in the following table. The definitions of these types are flexible to enable portability of code. Code requiring, e.g., a strict bit size should use inquiry functions to check that a given type meets the expected requirements.

int	*short integer*	*System defined standard integer*	
long	*long integer*	*Integer capable of counting elements in larges possible array*	
int8	*8-bit integer*	*OR smallest integer holding ≥2 decimal digits*	*OR same as **long***
int16	*16-bit integer*	*OR smallest integer holding ≥4 decimal digits*	*OR same as **int8***
int32	*32-bit integer*	*OR smallest integer holding ≥9 decimal digits*	*OR same as **int16***
int64	*64-bit integer*	*OR smallest integer holding ≥18 decimal digits*	*OR same as **int32***
int128	*128-bit integer*	*OR smallest integer holding ≥36 decimal digits*	*OR same as **int64***

Integer constants may be defined with respect to any base between 2 and 62 by specifying the base followed by **'r'** and then the digits, with **'a'..'z'** and **'A'..'Z'** representing 10 to 36 respectfully (case insensitive).

```
123         2r101011101   16rdeadbeef
```

By default integer constants yield values conforming to **int**. Long integer constants should append an **l** to the value. For other integer types, append an underscore and the corresponding (assumed) number of bits:

```
889874324897284746                    ! long
255_8                                 ! int8
2r1010101010101010101010101010_32   ! int32
```

PM also defines a set of real (floating point) types:

real	*Standard (system defined) floating point value*		
double	*Standard (system defined) double precision floating point value*		
real32	*32-bit floating point*	*OR real number with ≥1 decimal digits in mantissa*	*OR same as **real***
real64	*64-bit floating point*	*OR real number with ≥15 decimal digits in mantissa*	*OR same as **real32***
real128	*128- bit floating point*	*OR real number with ≥24 decimal digits in mantissa*	*OR same as **real64***

Real constants must contain a decimal point and may contain an exponent preceded by **'e'**.

```
-5.2        2e3            4.2e-20
```

By default real constants yield values conforming to type **real**. To specify a **double** constant, append a **d**. For other real types append and underscore and the assumed number of bits:

```
-5.2d      ! double precision
3.2e-3     ! real
1.2e-2_32  ! real32
```

Complex types are defined as follows:

cpx	*Complex number formed from **real** components*	
double_cpx	*Complex number formed from **double** components*	
cpx64	*64-bit complex number*	*OR complex number based on **real32** components*
cpx128	*128-bit complex number*	*OR complex number based on **real64** components*
cpx256	*256-bit complex number*	*OR complex number based on **real128** components*

Imaginary constants terminate with a letter '**i**' or '**j**'.

```
3.0i          ! cpx
1.0i_64       ! cpx64
-2.2-30di     ! double_cpx
```

The following intrinsic types define groups of numeric types:

- **any_int is int, long, int8, int16, int32, int128**
- **any_real is real, double, real32, real64, real128**
- **any_cpx is cpx, double_cpx, cpx64, cpx128, cpx256**
- **int_num includes any_int**
- **real_num includes any_int, any_real**
- **cpx_num includes any_int, any_real, any_cpx**
- **num includes cpx_num**
- **std_int is int, long**
- **std_real is real,double**
- **std_cpx is cpx, double_cpx**
- **std_num is std_int, std_real, std_cpx**

Numeric type balancing

Values of different numeric types may appear together in numeric expressions. The rules more mixed types are:

> The most general type is selected for the results and the least general value converted to that type. Generality is defined as follows (least..most):

> > **int long int8 int16 int32 int64 int128 real double real32 real64 real128 cpxdouble_cpxcpx64 cpx128 cpx256**

The **div** (integer divide) operator applies numeric balancing, and then selects the result type from the balanced type as follows:

int long int8 int16 int32 int64 int128	*as balanced type*
real double	Long
real32	int32
real64	int64
real128	int128

Non-numeric types

The **bool** type contains the logical values **true** and **false**

The **string** type contains values that are arbitrary length character strings. ***String constants*** are enclosed by double quotes '""' and may include any character other than '""'. They may not span more than one line.

```
"Hello World"
```

The **name** type contains values that are valid PM names. Name constants are formed by preceding a name by the dollar symbol:

```
$x
```

The **proc** type contains values that are valid PM procedure names (including operator symbol). Procedure constants are formed by preceding a name by the **proc** keyword:

```
proc cell_model
proc +
```

The **proc** keyword may be omitted for procedure names which are PM names (but not operator symbols or reserved words) which are not the same as the names of local variables, constants or references and not the same as any visible parameter definition.

The type **type** contains types used as first class values. Type values are created using:

> **type** *type*

Structures and records

Structures and records provide a mechanism to create aggregate values. They are very similar, except that it is not possible to alter a single component of a record, while this is possible for a structure.

A structure or record values is created using a generating expression:

> [**struct** | **rec**] [*name*] '{' *name* '='*exp* { ';' *name* '='*exp*} '}'

For example:

```
struct { name="John Smith", age=42 }
rec point { x=2, y=4 }
```

A structure or record value associates each component with a name. The structure of record value may optionally tagged with a further name. Component names are local to the structure or record. The name tagging the whole structure or record resides in a name space that is global to the program (unless preceded by an '_' in which case it is local to the given module)

A structure or record type is constructed using:

> [**struct** | **rec**] [*name*] '{'*name* [':' *type*] { ',' *name* [':' *type*] } '}'

For example:

```
struct { name:string, age:int }
rec point { x:int, y:int }
```

A structure or record value is a member of a given structure or record type if all of the following apply:

1. The value is a record and the type is defined using **rec** or the value is a structure and the type **struct**
2. The name tagging the value matches the corresponding name tag given in the **struct/rec** type, or both are absent
3. The same component names are present (not necessarily in the same order)
4. If a type constraint is associated with a component name in the type constraint expression then the corresponding component of the structure or record value, associated with the same name, conforms to that type constraint.

A structure or record type *T* conforms to a structure or record type *U* if:

1. Either both *T* and *U* are record types or both are structure types.
2. The name tag for *T* is the same as the name tag for *U*, or both are absent.
3. *T* and *U* have the same number of components with the same component names (not necessarily in the same order).
4. If for any given component name, *U* specifies a type constraint for that name, then *T* must also specify a constraint type conforming to the constraint type specified by *U*.

A given component of a structure or record may be accessed using the '**.**' operator. Structure components may be modified using the same operator.

```
person:=struct { name="John Smith", age=42 }
location:=rec point { x=2, y=4 }
person.age=person.age+1
print(person.age)
print(location.x)
```

In order to maintain the no-modification rule, a '**.**' operation by not be applied to a record on the left hand side of an assignment, on the right hand side of a reference connection for a non-constant reference or in the expression defining a reference parameter in a procedure call ('**&**')

Polymorphic types
A polymorphic value has an internal representation that is capable of representing any value in a given type. A value of a given polymorphic type is created using a generating expression:

'**<**' *type'* **>**' *subterm*

If a type is not given then it is set equal to the concrete type of the expression (see section on variables below). The value contained within a polymorphic value may be obtained using the unary '*******' operator.
For example:

```
Type int_or_string is int,string
a:= <int_or_string> 1
a= <int_or_string>'Hello'
print(*a)
```

A polymorphic type expression has the following form:

'**<**'[*type*]'**>**'

A polymorphic value conforms to type **<T>** if the value it contains conforms to *T*. A polymorphic type **<T>** conforms to **<U>** if *T* conforms to *U*.

Optional types
An optional type consists of the union of a given type and the **null** type – the latter is an intrinsic type containing the single value **null**.

opt *type*

There are no direct values of an optional type – they must be used in conjunction with polymorphic types or arrays.

Ranges

Range types represent open or closed intervals. Range values are created using the '**..**', '**<..**','**..<**' or '**<..<**' operators. For numerical values mixed type promotion rules apply. E.g:

```
closed_range:= 0..10
part_open_range:= 0<..10          ! Contains 1..10
part_open_range2:= 0..<10         ! Contains 0..9
real_open_range:=0<..<10.0        ! Contains { x | x>0.0 and x<10.0 }
```

The two bounds of a range value *R* may be obtained using the procedures **low(***R***)** and **high(***R***)**.

The following intrinsic range type are defined:

- **range_base includes any_real**
- **range { t: range_base }**
- **open_start_range { t: range_base }**
- **open_finish_range { t: range_base }**
- **open_range { t: range_base }**

Sequences

A sequence type represents a sequence of values. Sequence values are created by applying the **by** operator to a range value. The second argument to **by** gives the step between sequence elements. Numeric balancing will apply between the start/finish values of the range and the step argument.

```
integer_sequence := 1..6 by 2
real_sequence := 3.2 .. 5.4 by 0.7
```

The step of a sequence *S* may be found using the procedure **step(***S***)**. The two bounds of the underlying range value may be obtained using the procedures **low(***S***)** and **high(***S***)**.

Sequence types are defined as:

- **seq { t: range_base }**

Sets

Set types incorporate arbitrary sets of values. Set values are formed using a set generator expression:

> '**{**' [*exprlist*] '**}**' [**of** *type*]
> *exprlist ::=*
>   *expr* { '**,**' *expr* }

By default the element type for the set is equal to the union of the concrete types of all of the elements in the set. It is also possible to specify the element type using an **of** clause.
For example:

```
colours := { $red, $green, $blue, 16r1A3F2A }
numbers:= { 1 ,2, sqrt(2), a+b } of double
null_set_of_double = {} of real
```

The **in** operator allows set membership to be checked:

```
if colour in colours then ...
```

The intrinsic set type is given by:

- **set** { T }

A set value conforms to type **set{T}** if its associated element type conforms to *T*. A set type **set{T}** conforms to **set{U}** if *T* conforms to *U*.

Grids

A grid type defines an *N*-dimensional (2<*N*≤7) grid of points. Each dimension of the grid may be defined by a value conforming to the intrinsic type **grid_base**. By default this is either (1) an integer range (2) a numeric sequence (3) a set – in which case the ordering of elements along this dimension is arbitrary. There are seven specific grid types of varying number of dimensions, and a generalised grid type encompassing all of these:

- **grid_base includes range{std_int}, seq{}, set{}**
- **grid{T:grid_base} is T**
- **grid1d is grid{grid_base}**
- **grid2d is grid{grid_base,grid_base}**
- **grid3d is grid{grid_base,grid_base,grid_base }**
- **grid4d is grid{grid_base,grid_base,grid_base ,grid_base }**
- **grid5d is grid{grid_base,grid_base,grid_base ,grid_base ,grid_base }**
- **grid6d is grid{grid_base,grid_base,grid_base,grid_base,grid_base,grid_base}**
- **grid7d is grid{grid_base,grid_base,grid_base ,grid_base ,grid_base ,grid_base ,grid_base}**
- **grid is grid1d, grid2d, grid3d, grid4d, grid5d, grid6d, grid7d**

Grid values are created by calling the **grid** procedure:

```
model_grid := grid (0..num_cols, 0..num_rows)
```

The **grid** procedure will also accept integer arguments. These are converted to integer ranges starting from **1**.

The value corresponding to each element of a grid value *G* may be obtained using the procedures **d1**(*G*), **d2**(*G*), … , **d7**(*G*).

Domains

A domain defines a set of indices over which a (possibly parallel) iteration may be performed and over which arrays may be defined. The intrinsic domain type domain is defined as:

- **type dom includes grid_base, grid, vect_dom, mat_dom**

All valid domain values *D* support the **num_elem(D)** procedure which returns the number of elements in the domain.

Domain shapes

All domains *D* have a corresponding **domain shape**, which may be determined using the **shape(D)** procedure. This is defined as follows:

- The domain shape of a vector or matrix domain is equivalent to the original domain
- The domain shape of a grid domain **grid{$T_1,T_2,...$}** is given by
  proc shape(g: grid{$T_1,T_2,...$})= grid(1..num_elem(d1(g)),1..num_elem(d2(g), …)
- The domain shape of any other domain type *D* is given by **num_elem(D)**
- The programmer may provide specialised domain shape definitions for given types by providing additional definitions for the shape procedure. There is no restriction on the valid value types for a domain shape.

14

Arrays

An array type is a generic type indicating the storage of a set of values over corresponding to elements of a given domain.

One and two dimensional array values, with domains **int** and **grid{int,int}** respectively, may be defined using the following syntax:

> **['** *list2d* **']'**
> *list2d*::=
> > *exprlist* {**';'** *exprlist* } [**';'**]

For example:

```
array_1d := [ 1, 2, 3 ]                  ! domain is 1..3
array_2d := [ 1, 2, 3 ; 4, 5, 6; 7, 8, 9]   ! domain is grid(1..3,1..3)
```

For array values over other domains, or with elements which can store more generalised sets of values, the following extended syntax is available:

> **['** *list2d* **']'** [**of** *type*] [**over** *term*]

The **of** clause specifies a type constraint indicating the set of value that may possibly be stored in array elements. The **over** clause indicates the domain over which the array is defined. If the expression in the over clause yields an integer value this is converted to an integer range starting from **1**.

If the array definition contains a single value inside the square brackets, then this is replicated over all elements of the array. For a one dimensional list, the number of elements in the list must equal the number of elements in the specified domain. For a two-dimensional list of size *M* by *N*, the domain must be a two dimensional grid with shape **grid(***M***,***N***)**.

Array assignment requires that the two participating arrays have the same domain shape. Arrays cannot change size – flexible size arrays must be implemented using polymorphic types. An array of type **int[int]** cannot change size. A value of type **<int[int]>** can refer to arrays of different sizes.

Array values may also be created using the **dim** operator:

> *expr* **dim** *expr*

e.g.: `0.0 dim grid(0..3,0..3).`

This operator is defined as:

> **proc dim(element,domain:dom)= [element] over domain**

As with other PM operators, the **dim** operator may be defined for other or more specific types (see later).

Array types are defined using:

> [*type*] **'['** *opttypelist* **']'**
> *opttypelist*::=
> > [*type*] { **','** [*type*] }

An array value conforms to array type T[U] if its element type conform to T and its domain conforms to U. If more than one type is present within the square brackets then the domain constraint is equal to a grid type parameterised by the supplied type list. For example:

```
int[int]      ! Integer valued array over integer domain
int[]         ! Integer valued array over any domain
[]            ! Any array over any domain – equivalent to
int[int,int]  ! Integer array over gridded domain
              !– equivalent to int[grid{int,int}]
int[,,]       ! Integer array over any 3d gridded domain –
              !  -equivalent to int[grid{grid_base,grid_base,grid_base}]
```

Vectors and matrices

A vector is a numeric array defined over the **vect_dom** domain, which is a special one dimensional domain with a long integer dimension. A **vect_dom** value is created using the **vect_dom** procedure:

```
zero_vector = 0.0 dim vect_dom(1l..3l)
```

A matrix is a numeric array defined over the **mat_dom** domain, which is a special two dimensional domain with two long integer dimensions. A **mat_dom** value is created using the **mat_dom** procedure:

```
zero_matrix = 0.0 dim mat_dom(1l..3l,1l..3l)
```

Vector and matrix values may be created using a generating expression:

```
        '('list2d ')'
list2d::=
        exprlist {';' exprlist } [';']
```

For example:

```
        V :=   (1, 3, 1)
        a : = ( 3, 0, 0
               0, 2, 1
               0, 0, 1 )
```

Note the use of the optional semicolon rule in this example. The element type of a matrix or vector defined in this way is obtained by applying numeric balancing rules to the listed element values. There is no equivalent of the **of** or **over** clauses in a vector or matrix generator.

Matrix and vector values follow the usual rules for matrix multiplication. Vectors act as row or column matrices as appropriate in the context of matrix multiplication.

The vector and matrix intrinsic types are defined as:

- **vect is [vect_dom]**
- **mat is [mat_dom]**

Partitions and distributions

A ***partition*** of domain D over domain G associates a subdomain of D with each point in G. Conceptually, a partition acts as an array of sub-domain values although it is not defined as such.

A ***distribution*** extends this concept a stage further. A distribution of domain D over domain H associates a partition of D with each point in H. A distribution conceptually consists of an array of arrays of sub-domains.

The corresponding intrinsic types are **part{d:dom,g:dom}** and **distr{d:dom,h:dom}.**

Partitions and distributions are generated using intrinsic procedures:

For example:

```
domain:= grid(16,16)
part_domain:= grid(4,4)
partition:= block(domain,part_domain)
block:=partition[2,2]                      ! yields grid(5..8,5..8)
distribution:= block_cyclic(domain,partition,block=grid(2,2))
sub_block:=distribution[2,2][1]            ! yields grid(3..4,3..4)
```

Variables and constants and references

Objects are created and associated with names using a declaration statement:

---

statement::=
  const *assignlist*
  definition
  assignment
definition::=
  (*name*|'_') { ',' (*name*|'_') } '**:=**'*xexpr*
assignlist::=
  assignment { ',' *assignment* } *subexp*
assignment ::=
  lhs '**=**' *expr*
  lhs'**,**' *lhs*{ '**,**' *lhs* } '**=**' *call*
  name '**->**' *name qualifier*

For example

```
 x:= 4
 const message = "Hello World"
 r -> a.x
 const p -> a.x[3].u
```

An object has a ***concrete type*** that defines the set of values it is capable of storing. The concrete type of an object is determined by as follows:

1. If the initial value conforms to any of: **int long int8 int16 int32 int64 int128 real double real32 real64 real128 cpx double_cpx cpx64 cpx128 cpx256 bool string name proc** then that type is the concrete type of the object.
2. If the initial value is a structure or record, then the concrete type is a structure or record type with component type constraints given by the concrete types of the corresponding components, determined by applying these rules recursively.
3. If the initial value is a polymorphic value created using **poly** *T* {*expr*} then the concrete type is **poly{*T*}**

17

4. If the initial value is an array, then the concrete type is an array type with element and domain types equal to concrete types obtained recursively from corresponding components of the array value.

If the name is declared using '**:=**', values in the object may be changed later in the program. If it is declared using **const** then no changes to its value are permitted. The restriction for **const** names is enforced by forbidding them from appearing on the left hand side of an assignment statement or as a reference argument to a procedure. A name is defined until the end of the block of statements in which it is defined. Local (variable, constant and reference - defined later) names may not local shadow names defined in enclosing blocks or clash with procedure parameter names.

The values stored in variables may change over time. Values are changed in one of two ways: (i) using an assignment statement and (ii) by passing the variable as a reference argument to a procedure.

```
proc double(&x)=x*2
x := 4       ! Value stored in x is 4
x = x + 1    ! Value changed to 5
double(&x)   ! Value changed to 10
```

Some procedures return more than one value. In this case, multiple left hand sides are permitted:

assignment ::=
        lhs '**=**' *expr*
        lhs '**,**' *lhs* { '**,**' *lhs* } '**=**' *call*
        name '**->**' *name qualifier*
lhs::=
        name qualifier | '**_**'

qualifier::=
         { '**.**' *name* | '**[**' *exprlist*'**]**' | '**{**' *exprlist*'**}**' }

For example:

```
x,y,_, z = returns_four_args(a)
```

An underscore may be used as a place filler, allowing a given returned value to be ignored. It is also possible for the left hand side of an assignment to refer to a component of an object.

```
a.f = b
a[3,2]=c
```

If the types of the right and left hand side of an assignment differ, then a call is made to the procedure convert:

        lhs = convert(rhs,lhs)

This procedure, which may be extended through conventional proc definitions, is expected to convert the value rhs to the type of the value lhs.

References
A reference declaration assigns a name directly to an existing object, without creating a new object first. A reference is both created and changed using the same statement form:

assignment ::=
        name '**->**' *name qualifier*

For example:

```
r ->x
const s ->y.f
r ->a[3..9]
```

A reference declaration has the same scope as a variable or constant declaration. References can change type and are intrinsically polymorphic. A constant reference does not permit modification of the object to which it refers. Unlike a non-constant reference, however, it may refer to a component, or sub-component of a record.

Procedures
A procedure defines an operation on a set of objects, which may change some of their stored values. Procedures may also return one of more values. A procedure declaration defines the name of the procedure, parameters on which it operates and their associated type constraints and any values returned.

```
proc square(x) = x**2

proc calc_stats(data: float dim any) do
    ……
    result=mean,std_dev
endproc
```

Procedure names may be simple names (with or without a leading underscore) or may be one of the following operators:

+ - div mod * / ** == /= > >= and or // => dim = [] {} opt
.. <.. ..< <..< by

Note the absence of **<** or **<=** from this list – these operators are always defined relative to > or >= respectfully.

Each procedure is associated with a **signature** consisting of the following:

1. The name of the procedure
2. The number of values the procedure will return.
3. The number of arguments the procedure will accept
4. The type constrains on the procedure's parameters
5. Which parameters are declared to be reference parameters

Keyword parameters (defined below) are <u>not</u> part of the signature.

A procedure signature *P* is said to be **strictly more specific** than another procedure signature *Q* if:

1. The procedure names are the same.
2. The numbers of parameters are the same (allowing for variable argument definitions in either procedure).
3. All reference parameters occur in the same position.
4. The type constraint for each parameter of *P* conforms to the type constraint for the equivalent parameter of *Q*
5. Type constrains are not strictly equal to one another for at least one parameter position.

A procedure call invokes a procedure, supplying it with a list of values and/or objects to operate on. If the call is part of an assignment statement (discussed in detail below) then the call also provides a list of objects to receive any values generated.

```
u, v = myproc(x,&y,z*2)
```

If a procedure call is nested inside another procedure call, then the nested call is assumed to return a single value. Thus the following are equivalent:

```
y = f(g(x))

const _anon = g(x)
y = f( _anon )
```

Reference parameters, denoted by **&**, indicate that the procedure may change the value stored in corresponding object. Any procedure call must precede the corresponding argument with an **&**.

When encountering a procedure call, PM finds all procedures with signatures that *match* the call. A call matches a procedure signature if :

1. The procedure name is the same.
2. The number of arguments equals the number of parameters defined by the signature, or is greater than or equal to the number of parameters for a variable arguments signature.
3. All reference parameters are associated with a reference argument, starting with an &, in the same position.
4. The object supplied for each argument conforms to the type constraint for the corresponding parameter.

If more than one procedure matches a given call then PM will try to find a procedure that is strictly more specific than all the other candidate procedures. If no such procedure exists, then the call is *ambiguous* and the procedure selection is arbitrary. This condition may be checked for by the PM system and warnings issued.

Keyword arguments, if present, are not considered when finding a matching procedure. However, if the matching procedure will not accept a supplied keyword parameter then an error condition will arise.

A procedure may accept a variable number of arguments:

```
proc set_numbers(&dest:int[],arg:int...) do  ....
```

From the body of the procedure, the excess arguments are accessed using **arg[]** and **argc**. **argc** returns the number of arguments in positions corresponding to and following that of **arg…** as an **int** value. **arg[]** must be supplied an **int** value in the range **1..argc**, otherwise it will raise an error condition. It returns the value of the corresponding argument**.**

```
for i in 1..argc do
    dest[i] = arg[i]
endfor
```

It is also possible to pass optional arguments *en masse* to a procedure call:

```
proc process(a) do
    print("Value is"//a)
enproc
```

```
proc process(a,arg...) do
    process(a)
    process(arg...)
endproc

process_values(1,2,3,4)
```

Arguments passed in this way <u>do</u> form part of the signature of the call – they are used in the matching process and do not necessarily have to correspond to the variable argument portion of the called procedures parameter list.

Keyword parameters provide optional values to the procedure call. They are specified by providing a default value. Keyword parameters follow all other parameters and **arg...** (if present). Keyword parameters <u>do not</u> form a part of the signature of the procedure. An error results if the procedure selected by signature matching is not able to receive the keyword arguments provided in the call.

```
procrun_model(params : model_params, iterations = 1000,
    relaxation = 0.02) do
        .......
endproc

run_model(my_param_set, relaxation=0.042)
```

It is possible to define a procedure that takes additional unknown keyword parameters and passes these on to other procedures. This is achieved using the **key** keyword:

```
proc process_opts(&optarray,first_opt=.false,second_opt=33,key...) do
    ....
    process_other_opts(key...)
endproc
```

These additional keyword arguments may only be passed to another procedure – there is no facility to inspect them individually.

<u>Expressions</u>
An expression consists of a set of nested procedure calls. The usual infix notation is provided for common mathematical operations, but this is *syntactic sugar* for procedure calls.

Operation	Priority	Equivalent call	Description
x[a,b]	1	proc [] (x,a,b)	Array subscript
-x	2	proc - (x)	Minus
x**y	3	proc ** (x,y)	Power
x*y	4	proc * (x,y)	Multiplication (including matrix multiplication)
x/y	4	proc / (x,y)	Division (float result)
x div y	5	proc div (x,y)	Division (integer result)
x mod y	5	proc mod(x,y)	Remainder
x + y	6	proc + (x,y)	Addition
x - y	6	proc - (x,y)	Subtraction
x == y	7	proc == (x,y)	Equals
x /= y	7	proc /= (x,y)	Not equal
x > y	7	proc > (x,y)	Greater than
x >= y	7	proc>= (x,y)	Greater than or equal
x < y	7	proc> (y,x)	Less than
x <= y	7	proc>= (y,x)	Less than or equal
x and y	8	proc and (x,y)	Logical and

x or y	9	`proc or (x,y)`	Logical or
x // y	10	`proc // (x,y)`	String concatenation
x .. y	11	`proc .. (x,y)`	Range creation
x <.. y	11	`proc <.. (x,y)`	Range creation
x ..< y	11	`proc ..< (x,y)`	Range creation
x <..< y	11	`proc <..< (x,y)`	Range creation
x by y	11	`proc by (x,y)`	Sequence creation
x => y	12	`proc => (x,y)`	If x is true then y else null optional value
x || y	13	`proc || (x,y)`	If x is not null (or null optional value) then x else y
x => y || z	13	`proc||(proc=>(x,y), z)`	If x is true then y else z
x dim y	14	`proc dim(x,y)`	Array creation – spread copies of x over domain y

Sub-expressions

Sub-expressions may be separated from the main expression using **where**.

```
s = a * -exp (b/a) where a= c/sqrt(b)
```

There may be multiple values defined after a **where** keyword. These clauses may not refer to each other, but it is possible to follow with a second where statement and associated clauses:

```
a = x **2 / y**3 where x = s/p, y=s/q where s = sqrt(p**2 + q**2)
```

Sub-expressions may also be references:

```
a = a * 2 where a-> a[x1..x2 by x3,x1..y2 by y3]
```

It is also possible to include a check clause after an expression. This raises an error if the associated logical expression does not yield a **true** value:

```
x= my_sqrt_fn(x)  check x**2==old_x where old_x=x
```

Check expressions may optionally not be compiled or executed.

Subscripts
Subscript expressions have the following syntax:

> '**[**' *sexpr*{ '**,**' *sexpr* }'**]**' | '**{**' *sexpr* { '**,**' *sexpr* } '**}**' }
> *sexpr*::=
>   *expr*| [*expr*] ('..' | '<..' | '..<' | '<..<') [*expr*] [**by** [*expr*]]

Subscript expressions using square brackets **[]** index an element or slice with reference to the domain of the array value being subscripted.   Subscript expressions using braces **{}** index elements or slices with respect to the shape of the domain of the array.

Subscript expressions may contain the keywords **high**, **low**, or **step**. If this occurs in a given index position of the subscript expression list, then these are converted to calls to the equivalent procedures on the corresponding dimensions of the array's domain.

```
x[low..high/2,low..high by step*2 ]
```

is equivalent to:

```
x[low(d1(x))..high(d1(x))/2,low(d1(x))..high(d1(x))/2 by step(d2(x))*2]
```

If the subscript expression consists of a range generator or a direct combination of range and sequence generators (**x..y by z**) then the start, end or step expressions may be omitted, in which case they are set to **low**, **high** or **step** respectfully. Thus the above expression could be written:

```
x [ .. high/2 , .. by step*2 ]
```

Statements
PM provides a range of conventional control statements:

> **if** *xexpr* **then** *statements* { **elseif** *xexpr* **then** *statements* } [**else** *statements*] **endif**
> **select** *xexpr* {(**case** *xexprlist* | **default**) **do** *statements* } **endselect**
> **while** *xexpr* **do** *statements* **endwhile**
> **repeat** *statements* **until** *xexpr*

The **if** statement conditionally executes a list of statements:

```
if x > y then
    print("X is greater")
endif

if x>y then
    print("X is greater")
else
    print("Y is greater or equal")
endif
```

The **select** statement tests an expression against lists of values and executes the first statement list to be associated with a value matching the expression. The **select** statement builds sets of values (which do not have to be constants) and uses the set **in** operator to check the applicability of each case.

```
select digit
    case '1','3','5','7','9' do print("Odd")
    case '2','4','6','8' do print("Even")
    otherwise do print("?")
endselect
```

The **repeat** statement sequentially repeats a list of statements until an expression yields a true value (test at the end).

```
repeat
    x = x/2
    y = y+1
until x == 0
```

The **while** statement executes a statement zero or more time while a given expression yields a true value (test at the start)

```
nbits=0
while x>0 do
    x = x/2
    nbits = nbits + 1
endwhile
```

In each of these cases, conditional expressions must return a **bool** value. An error will result if they do not.

Two statements allow for debugging and may optionally not be compiled and executed. The **check** statement confirms that an expression yields a **true** value and raises an error if this is not the case. The **debug** statement executes a block of code only if debugging is enabled.

```
        check  xexpr
        debug  statements enddebug
```

For example:

```
        check value /= null

        debug print("Entering main loop")enddebug
```

Parallel execution

The **for** statement executes its body in parallel for every element of one or more domains or arrays:

```
        for iter subexp [ seq|[ using assignlist] ] loopbody endfor
iter::=
        name { ',' name } in exp { ','exp }
        name from exp follow qualifier [ when exp ]
loopbody::=
        do statements
        find statements [ otherwise statements ]
```

For example:

```
        for pixel in image do
             pixel = min(pixel,threshold)
        endfor
```

It is possible to iterate over more than one domain and/or array, providing they all share the same domain shape:

```
        for pixel in image1, pixel2 in image2 do
            if pixel+pixel2>threshold then
                 pixel=pixel-threshold/2
                 pixel2=pixel2-threshold/2
            endif
        endfor
```

Statements in the body of a parallel loop are not normally allowed to modify variables defined outside of the scope of the **for** statement. It is possible to require that the loop body is executed sequentially, in which case these restrictions do not apply.

```
        for s in array seq do
             sum = sum + f(s)
        endfor
```

<u>Communicating operators</u>

> '**@**'*name*
> '**@**' *name* '**(**'*arglist*'**)**'
> *name* '**@**' '**{**' *sexpr* { '**,**' *sexpr* } '**}**'
> *name* '**::**' *name* ['**@**' '**{**' *sexpr* { '**,**' *sexpr* } '**}**']

Numerical models will usually require interaction between loop invocations. PM enables this using the '@' operator. Used as a unary operator (@x) this provides an array view of a given expression across all invocations. Used as a subscripting operator (x@{ndb_desc}) it returns values from a local neighbourhood of the calling invocation. For example:

```
for cell in model_array do
    advection=advection_model(cell)
    advected_cell=cell@{advection}
    cell=model(advected_cell,parameters)
endfor
```

In the latter case, the result will be an array with an optional element type and off-edge values undefined.

```
for pixel in image do
    pixel = median(pixel@{-1..1,-1..1})
endfor
```

The neighbourhood is defined with respect to the common shape of the domains/arrays being iterated over, shifted so that the zero point lies at the location of the point associated with the current invocation.

The '::' operator repeatedly applies a given binary operation applied between the value of a given expression across all invocations of the enclosing parallel loop:

```
for cell in model_grid do
    cell=model(cell,parameters)
    cell=cell/sum::cell     ! Normalise
enddo
```

A reduction operation must be defined using a specialised procedure definition:

```
proc sum:: (x,y) = x + y
```

The binary operation is applied with arbitrary ordering and bracketing. This will only give a deterministic result if the operation is both commutative and associative. A reduction procedure definition must specify only two parameters. These must have identical type constraints, if present.

It is possible to confine a reduction operator to a neighbourhood:

```
local_mean = sum::A@{nbd}/count:: A@{nbd}
```

Communicating procedures

It is possible to define **communicating procedures** incorporating '**@**' and '**::**'operators that can only be called inside parallel for statements:

```
proc @mean_filter(x,n)=sum(x@{nbd})/count(x@{nbd})
     where nbd=grid(-n..n,-n..n)

for pixel in image do
     pixel = @mean_filter(x,1)
endfor
```

Synchronisation requirements – structured parallelism
PM parallel communication operators must operate in the same sequence for each invocation of the enclosing parallel statement. A simplified set of rules enable static analysis of this communication matching requirement.

1. The sequence associated with a communication operator is defined to be the operator type (unary or binary '**@**', general or neighbourhood reduction) combined with the variable to which the operator is applied (in terms of the object accessed – not just the variable name)
2. The sequence associated with a call to a loop procedure is defined to be procedure body invocated.
3. The sequence associated with a statement list is defined to be an ordered set of the sequences associated with each statement in the statement list which contains communicating operators
4. For a conditional statement (**if**, **select**) the sequences associated with each branch of the conditional must be the same. The sequence for the statement as a whole consists of the sequences for any conditional/selecting expression, in textual order, followed by the sequence for an arbitrary branch of the conditional.
5. For a sequential loop (**while**, **repeat**, **for**..**seq**) the sequence is defined as the sequence of the loop body, grouped as a *repeated* sequence. The form of loop is not relevant; all sequential loops apply the same repeated grouping. In terms of sequence matching, a repeated group must match another repeated group containing the same sequence.

If a conditional expression (if, select) contains communicating operators then each conditional branch of the statement must contain the same sequence of operations.

If communicating operators appear within non-parallel loops (**while**, **repeat**, **for**..**seq**). These loops are not forced to repeat the same number of times for each invocation of the enclosing parallel **for**. If a communicating operator executed in a longer-running loop tries to access a value from an invocation where the loop has already terminated, the operator will be supplied with the value that the named variable possessed immediately prior to the termination of the shorter-running loop. There is an implicit synchronisation at the end of sequential loops – a communicating operator executing after the end of the loop in one parallel invocation cannot match an operator executed from within the loop in a different parallel invocation.

If a sequential loop containing communicating operators occurs in a conditional statement, then a loop containing the same sequence of communicating operators must occur at the same position (in terms of sequence of communicating operators) on each branch of the conditional statement. Matching loops do not have to employ the same statement forms – a **while** loop can match a **repeat** loop and both can match a **for**.. **seq** loop. Moreover, a nested set of sequential loops in which only the innermost loop body contains communicating operators may be considered to be a single loop when matching parallel communication operations.

The **for** .. **do** .. **endfor** version of the **for** statement described above always invokes the loop body for every element of the domain. A second version executes arbitrary invocations of the loop body until at least one such body encounters a **found** statement:

---

      found *assignlist*

Out of the loop invocations that encounter the **found** statement, only one will execute that statement. If all loop invocations complete without encountering a **found** statement, then the statements in the **otherwise** clause (if present) are executed. This clause is only executed once – not as a parallel invocation – and is thus not restricted with respect to side effects.

```
for pixel in image find
     if pixel>threshold then
            found high_pixel=pixel
     endif
otherwise
     high_pixel=-1
endfor
```

The body of a **find** loop may not use the **@** or **::** operators, except within nested parallel **for** – **do** statements.

Task parallelism

Task parallelism is supported by the **do**..**enddo** construct.

      [**using** *keyargs*] [**with** *statements*] **do** *statements* { **also do** *statements* } **enddo**

The semantics of a **do** statement are defined according to an equivalent **for** statement.

using *opts*   **with** *statements*$_0$ **do** *statements*$_1$ **also do** *statements*$_2$ … **also do** *statements*$_N$**enddo**

is semantically equivalent to:

> **for** *_index* **in 1..**N **using** *opts* **do**
>   *statements*$_0$
>   **select** *_index*
>   **case 1 do** *statements*$_1$
>   **case 2 do** *statements*$_2$
>   …
>   **case** N **do** *statements*$_N$
>   **endselect**
> **endfor**

Tasks and processors

PM programs are assumed to be executed on an N-dimensional array of ***processors***. Processors are capable of efficiently executing a set of ***tasks***. The PM specification does not define how the abstract concept of a processor into hardware: a processor can relate to anything from a single core to a sub-cluster. The primary requirement is that a processor is capable of keeping its own hardware busy running available tasks which effectively requires efficient mechanisms for task migration within the processor. Task migration between processors, however, is assumed to either introduce a degree of overhead which the programmer has to take into account, or not to be possible at all in any acceptably efficient manner.

PM assumes that the number of processors in the processor grid is fixed by the implementation. When a PM program proceeds, this processor set is divided among available processors according to the following:

> On start-up, the PM program is assumed to be running a single task. This task owns all processors in the grid.
>
> When a parallel statement (**for** or **do**.. **also do**) is encountered. Then processor allocation is governed by comparing its iteration space (which is implicit in **do**..**also do**) to the processor grid associated with the currently running task
>
> 1. If the number of elements in the processor grid is greater than the number of elements in the iteration space, then the processor grid is partitioned over the iteration space using the ***default processor partition function***. A new task is created to process each element in the iteration space. This task owns the processors listed in the element of the processor grid partition corresponding to the given element in the iteration space.
>
> 2. If the number of elements in the processor grid is less than or equal to the number of elements in the iteration space, then a distribution of the iteration space over the processor grid is calculated using the ***default iteration space distribution function***. The distribution is conceptually a nested array **D[p][i]** of subdomains of the iteration space, indexed by both an element **p** of the processor grid and a second, arbitrary, index **i**. On each processor, a new task is created for each element **D[p][i]** of **D[p]**. Each task is responsible for processing the iteration space elements contained in **D[p][i]** and shares ownership of the single-element processor grid **p.** The default distribution function applies a block cyclic distribution for matrix domains and a nested block distribution for grid domains.

The PM definition does not specify how a task owning multiple processors utilises those processors to perform its computations. Two possible implementation approaches are:

1. Every processor in the processor set owned by a given task redundantly executes the same code.

2. A single processor in the processor set executes the code, the remaining processors are held in a standby state until needed by nested parallel operations.

There is clearly a trade-off between ease of implementation, inter-processor message overhead and resource utilisation between these two approaches. A PM implementation is not required to employ either of these extremes. The only requirement, as with defining processors, is that controlling implementation overheads should not require intervention from the programmer.

The above process may be modified by providing keyword arguments in the using clause of a **for** or **do** statement. The following options are defined.

distr=_Distribution_	Use the given distribution in place of the output of the default distribution procedure
proc_distr=_Distribution_	_Use the given distribution in place of the output of the default processor distribution procedure_
work=_Array_	_If this option is specified, the default processor distribution will unevenly distribute processors according to a measure of the relative amount of work associated with each invocation. The given array should have numeric elements and a domain shape equal to the common shape of the inputs to the parallel statement_

float=_BoolValue_   _If this option is true, then the parallel statement should allow tasks to move between processors, assuming the PM implementation supports this option. The default is false._

block=_Domain_   _This option sets the block size used by the default distribution procedure._

Syntax
module::=*program_module* | *library_module*
program_module::=
    [*decls*] *statements*
library_module::=
    decls [**debug** *statements* **enddebug**]
decls::=
    { *import* '**;**' }(*import*| *decl*) { '**;**'*decl* }
import::=
    use *modname* [*modifiers*] { '**,**' *modname* [*modifiers*] }
modname::=
    name{ '**.**' *name* }
modifiers::=
    '**{**' *modifier* { '**,**' *modifier* } '**}**'
modifier::=
    [**type** | **param** | **proc**] *name* '**=>**' *name*
decl::=
    procdecl | *typedecl* | *paramdecl*
signature::=
    procsig | *typesig* | *paramsig*
procdecl::=
    proc *procname* *params*'**=**' *xexprlist* [**do** *statements* **endproc**]
    proc *procname* *params* [**check** *exprlist*] **do** *statements* [**result** '**=**'*xexprlist*] **endproc**
procname::=
    [*name* | *op*| *name* '**::**'| **low** | **high** | **step** | '**@**' *name*]
op::=
    '**+**' | '**-**' |**div** | **mod** |'*****' | '**/**' | '******'| '**==**' | '**/=**' | '**>**'| '**>=**' | **and** | **or** | '**//**' | '**=>**' | '**||**' | **dim**| '**=**'| '**[]**' | '**{}**'| **opt**|
    '**..**' | '**<..**' | '**..<**' | '**<..<**' | **by**
params::=
    '**(**'{ *param* '**,**' } [*param*['**,**' *keyparams*] | **arg** '**...**' ['**,**' *keyparams*] | *keyparams*] '**)**'
param::=
    ['**&**'] *name*['**:**' *type*]
keyparams::=
    { *name* '**=**' *expr*'**,**'} (*name* '**=**' *expr* |**key** '**...**')
namelist::=
    name { '**,**' *name* }
typedecl::=
    type *name* ['**{**' *typeparams* '**}**'][**in** *namelist*] (**is** *typelist* | **includes** *typelist* | **also includes** *typelist*)
typeparams::=
    name '**:**' *type* { '**,**' *name* '**:**' *type* }
paramdec::=
    param *name* '**=**' *xexpr*
typelist::=
    type { '**,**' *type* }
opttypelist::=
    [*type*] { '**,**' [*type*] }
type::=
    name ['**{**' *opttypelist* '**}**']
    [*type*] '**[**' *opttypelist* '**]**'
    '**<**' *type* '**>**'
    opt *type*
    [**struct**| **rec**] [*name*] '**{**'*name* ['**:**' *type*] { '**,**' *name* ['**:**' *type*] } '**}**'
statements::=
    statement{'**;**'*statement*} ['**;**']

statement::=

      If *xexpr* **then** *statements*{ **elseif** *xexpr* **then** *statements* } [**else** *statements*] **endif**

      select *xexpr* { (**case** *xexprlist* | **default**) **do** *statements* } **endselect**

      while *xexpr* **do** *statements* **endwhile**

      repeat *statements* **until** *xexpr*

      for *iter subexp* [**seq**|[**using** *keyargs*]] *loopbody* **endfor**

      [**using** *keyargs*] [**with** *statements*] **do** *statements* { **also do** *statements* } **enddo**

      check *xexpr*

      debug *statements* **enddebug**

      found *assignlist*

      const *assignlist*

      let *assignlist*

      assignment

      call

iter::=

      name{ '**,**' *name* } **in** *exp* { '**,**'*exp*}

      name **from** *expr* **follow** *qualifier* [**when** *expr*]

loopbody::=

      do *statements*

      find *statements* [**otherwise** *statements*]

definition::=

      (*name*|'**_**') { '**,**' (*name*|'**_**') } '**:=**'*xexpr*

assignlist::=

      assignment { '**,**' *assignment* } *subexp*

assignment ::=

      lhs '**=**' *expr*

      lhs'**,**' *lhs*{ '**,**' *lhs* } '**=**' *call*

      name '**->**' *name qualifier*

lhs::=

      name qualifier |'**_**'

expr::=

      <u>*lowest to highest precedence*</u>

      expr **dim** *expr*

      expr '**||**' *expr*

      expr '**=>**' *expr*

      expr['**..**' | '**<..**' | '**..<**' | '**<..<**'| **by**]*expr*

      expr'**//**'*expr*

      expr **or** *expr*

      expr **and** *expr*

      not *expr*

      expr ['**==**' | '**/=**' | '**>**'| '**<**' | '**>=**' | '**<=**' | **in**] *expr*

      expr ['**+**' |'**-**'] *expr*

      expr **mod** *expr*

      expr **div** *expr*

      expr['*****'|'**/**'] *expr*

      expr '******' *expr*

      '**-**'*term*

      term qualifier

subexp::=

      [**check** *exprlist*] *whereclause*

whereclause::=

      { **where** *namelist* '**=**'*expr* { '**,**' *namelist* '**=**'*expr* } }

xexpr::=

      expr subexp

xexprlist::=
        exprlist subexp
exprlist ::=
        expr { ',' *expr* }
call::=
        [*name*| **low** | **high** | **step**]'**(**'*arglist*'**)**'
        proc *procname* '**(**'*arglist*'**)**'
term::=
        subterm
        value
        '**@**'*name*
        '**@**' *name* '**(**'*arglist*'**)**'
        name '**@**' '**{**' *sexpr* { ',' *sexpr* } '**}**'
        name '::' *name* ['**@**' '**{**' *sexpr* { ',' *sexpr* } '**}**']
subterm::=
        '**<**' [*type*] '**>**' *subterm*
        name
        '**(**'*expr*'**)**'
        call
        array
        [**struct** | **rec**] [*name*] '**{**' *name* '='*exp* { '**;**' *name* '='*exp* } '**}**'
value::=
        constant
        arg'['*expr*']' | **argc** | **low** | **high** |**true** | **false**
        proc *procname* | **type** *type*
constant::=
        number| *string* |'**$**' *name*
arglist ::=
        { *arg* ',' } (*arg* | **arg** '**...**')[',' *keyargs*]
        [*keyargs*]
arg::=
        '**&**'*name qualifier* | *expr*
keyargs::=
        { *name* '=' *expr* ',' } (*name* '=' *expr* | **key** '**...**')
qualifier::=
        { '**.**'*name*| '**[**' *sexpr* { ',' *sexpr* } '**]**' | '**{**' *sexpr* { ',' *sexpr* } '**}**' }
sexpr::=
        expr| [*expr*] ('**..**' | '**<..**' | '**..<**' | '**<..<**') [*expr*] [**by** [*expr*]]
array::=
        ['**(**' (*list2d* | generator) '**)**'
        '**{**' [*exprlist* | *generator*] '**}**' [**of** type]
        '**[**' (*generator* | list2d) '**]**' [**of** *type*] [**over** *term*]
list2d::=
        exprlist {'**;**' *exprlist*} ['**;**']
generator::=
        expr '**:**' *iter* { '**;**' *iter* } ['**|**' *expr*]

<u>Intrinsic procedures</u>

FORTHCOMING